

Legitimate Intervention Framework (LIF)

Contribution to GnosisDAO's 'Framework for the Future' Consultation

Elem Oghenekaro

January 2026

Contents

0.1	Executive Summary	2
0.2	Part 1: Problem Space	2
0.2.1	1.1 Scale of the Threat	2
0.2.2	1.2 Threats from AI Acceleration	2
0.2.3	1.3 New Reality	2
0.3	Part 2: The Hierarchy of Precision	3
0.3.1	Principle of Minimal Intervention	3
0.4	Part 3: Legitimacy Conditions	3
0.4.1	Legitimacy Test	4
0.4.2	Decision Parameters	4
0.5	Part 4: The Procedural Framework	5
0.5.1	4.1 Pre-Incident Phase	5
0.5.2	4.2 Incident Response Phase	5
0.5.3	4.3 Post-Incident Phase	6
0.6	Part 5: Addressing the Colonial Critique	6
0.6.1	Response:	6
0.7	Part 5.5: Organisational Resilience	7
0.7.1	The Resilience Stack	7
0.7.2	High-Reliability Industry Principles	7
0.7.3	The Accountant's Question	8
0.7.4	Integration with the Framework	8
0.8	Part 6: Alignment with Gnosis Values	8
0.9	Proposed Next Steps	8
0.9.1	For GnosisDAO	8
0.10	Data Sources & References	9
0.11	Appendix: Glossary	9

Response to Sebastian Bürgel (@SCBuergel)'s "A Framework for the Future" consultation

This document synthesises findings from the Legitimate Intervention Framework (LIF) (*formerly Native Compliance Framework (NCF)*) initiative to inform GnosisDAO's development of intervention criteria for its emergency framework. It offers a structured approach to the "credible neutrality vs. user safety" tension identified in the consultation.

0.1 Executive Summary

The Balancer exploit (November 2025) and Gnosis Chain's response raised the need to define when and how intervention should occur: intervention capabilities exist, but no shared criteria exist for when to use them.

This framework proposes:

1. A Hierarchy of Precision for intervention scope (from “Nuclear Option” to surgical)
2. Legitimacy Conditions that must be satisfied for any intervention, and
3. A Procedural Framework aligned with GnosisDAO’s governance values

Key Claim: Legitimate intervention is not the opposite of decentralisation but a formalisation of emergency powers that already exist but remain opaque or unclear, as a way of avoiding improvisational responses.

0.2 Part 1: Problem Space

0.2.1 1.1 Scale of the Threat

Research from Charoenwong & Bernardi (2022) documents \$88 billion in cryptocurrency losses (2011–2021), with the following distribution:

Category	Incidents	%
Security Breach	20	66%
Human Error	5	17%
Insider Theft	5	17%

Post-2021 Update: The LIF database now tracks 759 exploits totalling \$91.3 billion (2014–2025).

0.2.2 1.2 Threats from AI Acceleration

The threat landscape has fundamentally shifted. Anthropic’s December 2025 Red Team report reveals:

- AI agents can autonomously execute 55.8% of 2025 exploits
- Exploit revenue potential doubles every 1.3 months
- Cost per contract vulnerability scan: \$1.22

The implication is that manual governance processes cannot keep pace with AI-accelerated attacks. Pre-defined intervention criteria that can be invoked rapidly are essential.

0.2.3 1.3 New Reality

A November 2025 analysis by Bybit’s Security Lab examined 166 blockchains:

Capability	Chains	%
Confirmed Freezing	16	9.6%
Potential Freezing	19	11.4%
No Freezing	131	79%

Key Finding: 21% of the market already has intervention capability. The question is not *whether* to intervene, but *how to do so legitimately*.

0.3 Part 2: The Hierarchy of Precision

Current debate tends toward binary thinking (Halt vs. Nothing). This framework proposes a **graduated response spectrum**:

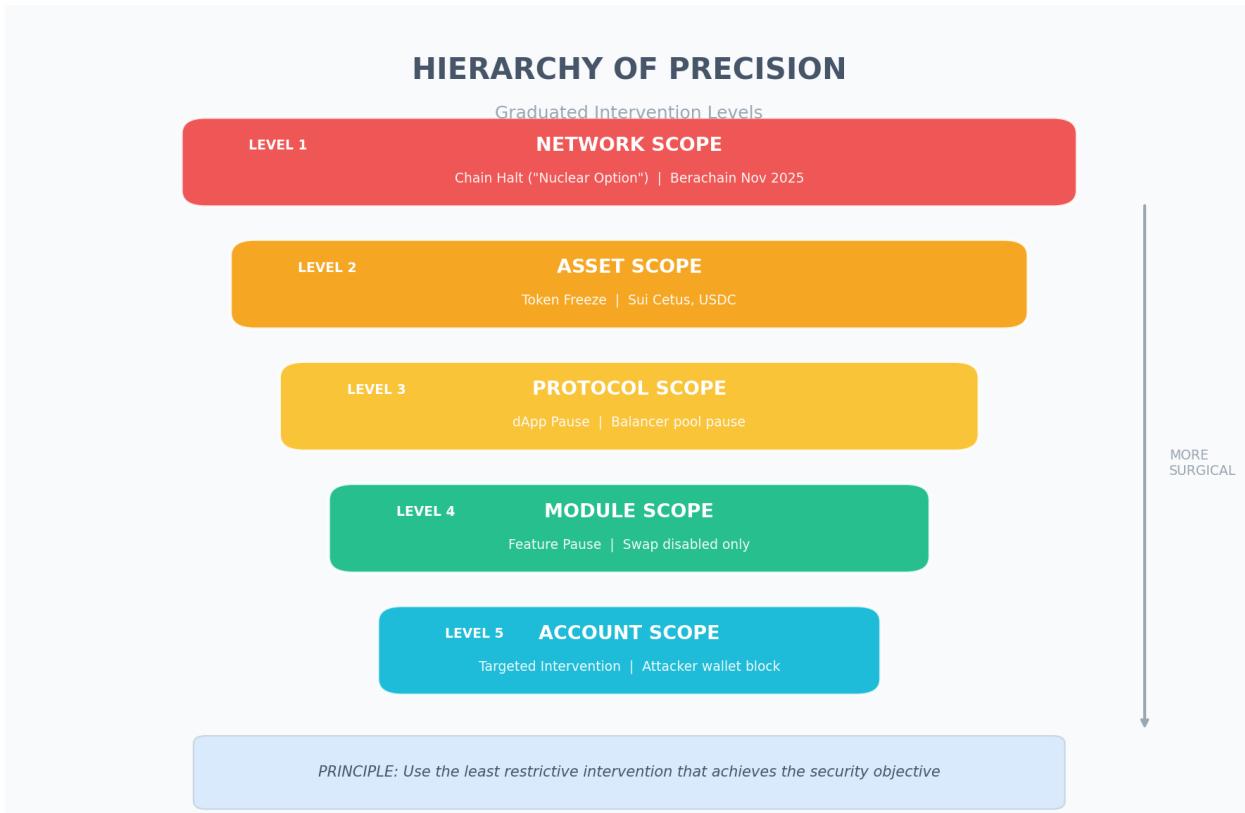


Figure 1: The Hierarchy of Precision — graduated intervention levels from Network Scope (most disruptive) to Account Scope (most surgical).

0.3.1 Principle of Minimal Intervention

The least restrictive intervention that achieves the security objective should be preferred.

Example Analysis (Balancer Nov 2025):

- **Berachain** used Level 1 (Chain Halt) — effective but maximally disruptive
- **Sonic** used Level 5 (Account Freeze) — surgical but required rapid deployment
- **Optimal:** Pre-defined Level 4/5 capabilities with clear activation criteria

0.4 Part 3: Legitimacy Conditions

Drawing on Charoenwong et al.'s “A Framework for Evaluating Regulatory Approaches to On-Chain Decentralised Software” (2025), five regulatory models can be identified:

Model	Description	Authority
Centralised	Direct government enforcement	Human
Gatekeeper	Delegated to intermediaries	Human
RegTech	Algorithmic monitoring	Algorithmic + Human

Model	Description	Authority
Ex Ante Automated	Rules enforce before execution	Algorithmic
Ex Post DeFi	Community governance after incident	Human

For Gnosis, a hybrid of Models 3–4 is recommended: algorithmic detection with human governance for activation. This approach balances the speed and precision of automated systems with the accountability and legitimacy of human oversight.

0.4.1 Legitimacy Test

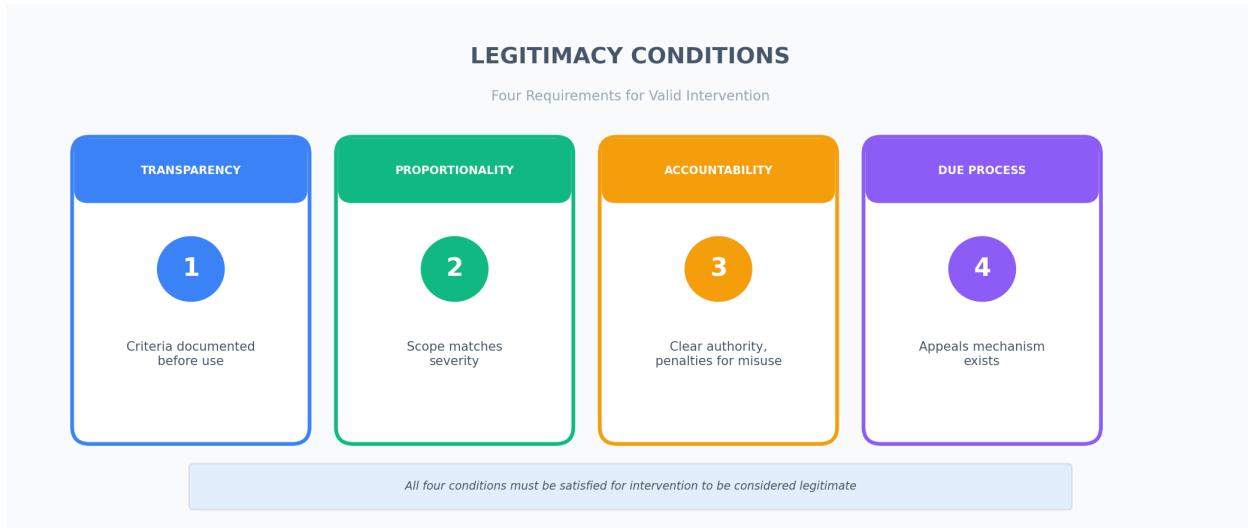


Figure 2: The four legitimacy conditions that must be satisfied for intervention to be considered valid.

For any intervention to be considered legitimate, it should satisfy:

1. Transparency Condition
 - Intervention criteria are publicly documented *before* use
 - Decision-making process is visible to the community
 - Technical implementation is auditable
2. Proportionality Condition
 - Scope matches severity (Hierarchy of Precision)
 - Time-limited by default (sunset clauses)
 - Reversibility preserved where possible
3. Accountability Condition
 - Clear authority chain (who can trigger?)
 - Post-incident review required
 - Penalty for misuse (slashing, removal)
4. Due Process Condition
 - Affected parties can contest (within constraints)
 - Appeals mechanism exists
 - Evidence standards defined

0.4.2 Decision Parameters

In the Gnosis Community AMA on 7 January 2026, Gnosis cofounder Friederike Ernst outlined specific factors the core team is considering for intervention decisions. Incorporating these parameters into the framework is strongly supported:

Parameter	Description	Relevance
Exploit Type	What kind of attack occurred (reentrancy, oracle manipulation, access control, etc.)	Determines appropriate intervention level
Protocol Type	DeFi, bridge, NFT, infrastructure	Different risk profiles require different responses
Exploit Novelty	Known vulnerability vs. zero-day	Novel exploits may warrant faster response
Security Claims	What security guarantees did the protocol advertise?	Affects accountability assessment
Audit Status	Was the protocol audited? By whom? When?	Informs proportionality of response
Fund Percentage	What percentage of Gnosis Chain TVL is affected?	Higher impact justifies broader intervention

These parameters should inform both the *trigger threshold* (when to intervene) and the *intervention level* (how broadly to intervene). The following are recommended:

1. Pre-defined thresholds for each parameter documented in the framework
2. Weighted scoring to combine multiple factors into a decision matrix
3. Transparent criteria published before any incident occurs

0.5 Part 4: The Procedural Framework

0.5.1 4.1 Pre-Incident Phase

Establish the Rules Before the Crisis

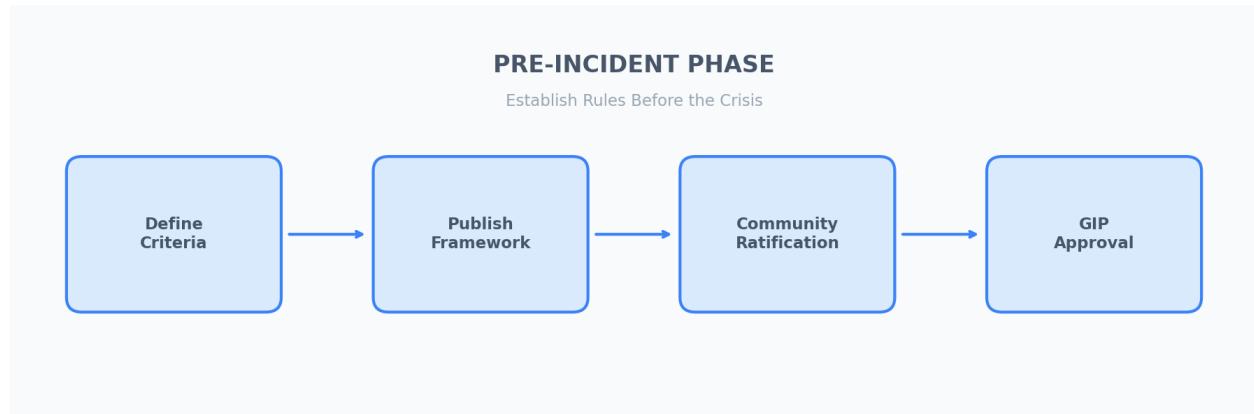


Figure 3: The pre-incident phase establishes intervention criteria through community ratification.

Required Documentation: - [] Intervention type definitions (Levels 1–5) - [] Trigger conditions for each level - [] Authority designation (multisig? quorum? node operators?) - [] Time limits and sunset clauses - [] Review and appeal procedures

0.5.2 4.2 Incident Response Phase

An “Optimistic Freeze” Model

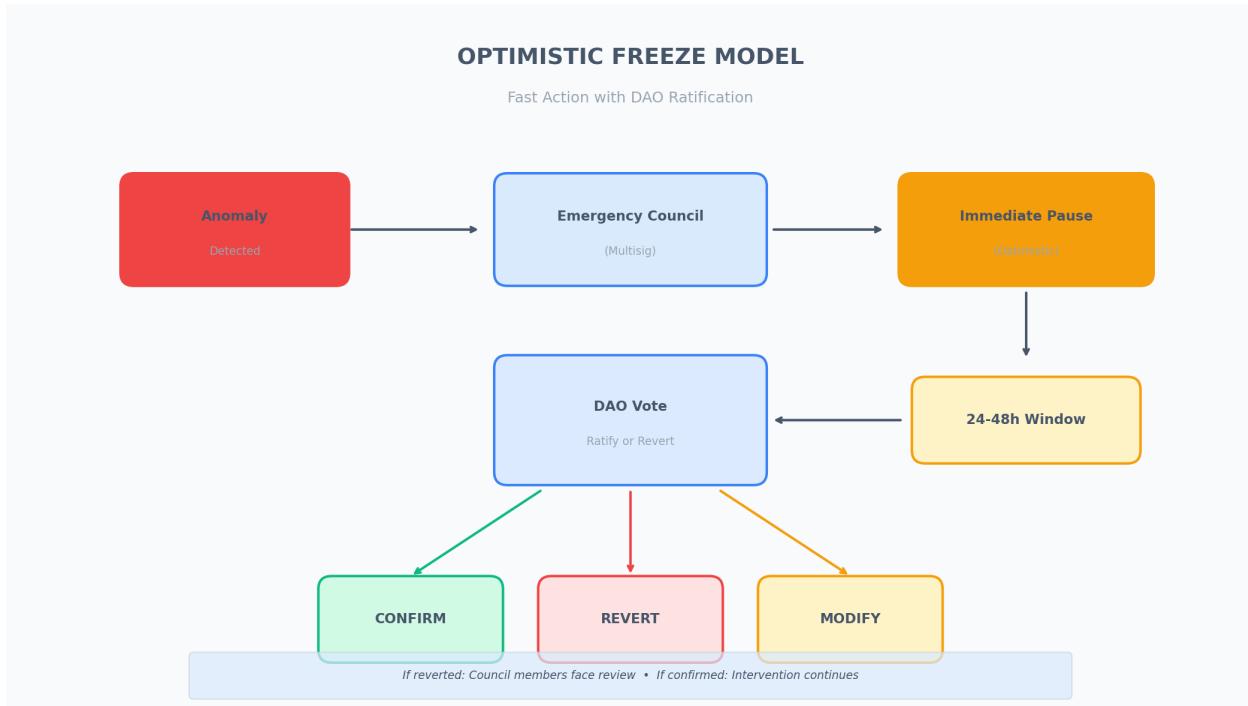


Figure 4: Optimistic Freeze model enables fast action with mandatory DAO ratification.

Key Features: - **Speed:** Emergency Council can act in minutes - **Legitimacy:** DAO ratification required within 24–48h - **Accountability:** If reverted, Council members may be subject to review

0.5.3 4.3 Post-Incident Phase

Mandatory Transparency

Every intervention requires: 1. Public Post-Mortem (within 7 days) - What happened - What intervention was taken - Why that level was chosen - Outcome and funds status

2. Community Review Period (14 days)
 - Open discussion on forum
 - Assessment of whether criteria were correctly applied
 - Proposals for framework amendments
3. Lessons Repository
 - Documented for future reference
 - Informs framework updates

0.6 Part 5: Addressing the Colonial Critique

@mrtdlgc raised a valid concern about “legitimising censorship.” This is taken seriously.

0.6.1 Response:

1. Transparency is the Antidote to Colonialism
 - Colonial systems operate through hidden power structures
 - This proposal makes intervention criteria public and contestable
 - The community decides the rules rather than a select few
2. The Alternative is Worse

- 21% of chains already have freezing capability, mostly undocumented
 - Without a framework, intervention happens anyway but without accountability
 - Gnosis is pioneering the path for a formalised system the industry can adopt or adapt
3. The Hierarchy Preserves Sovereignty
 - Level 5 (Account Scope) affects only proven attackers
 - Level 1 (Network Scope) requires highest justification
 - Proportionality limits scope by design
 4. Community Governance Enables Accountability
 - GIP process = collective decision-making
 - Node operators can refuse to implement disputed changes
 - Fork remains the ultimate exit right
-

0.7 Part 5.5: Organisational Resilience

Another important dimension is *human factors*. A framework is only as strong as the people executing it. This section draws from @mfw78's insights on aviation, nuclear power, and emergency medicine:

"Improvisation is often a symptom of under-resourcing." — @mfw78

If GnosisDAO determines that intervention capability is appropriate, that capability needs to be *designed, resourced, and sustained*.

0.7.1 The Resilience Stack

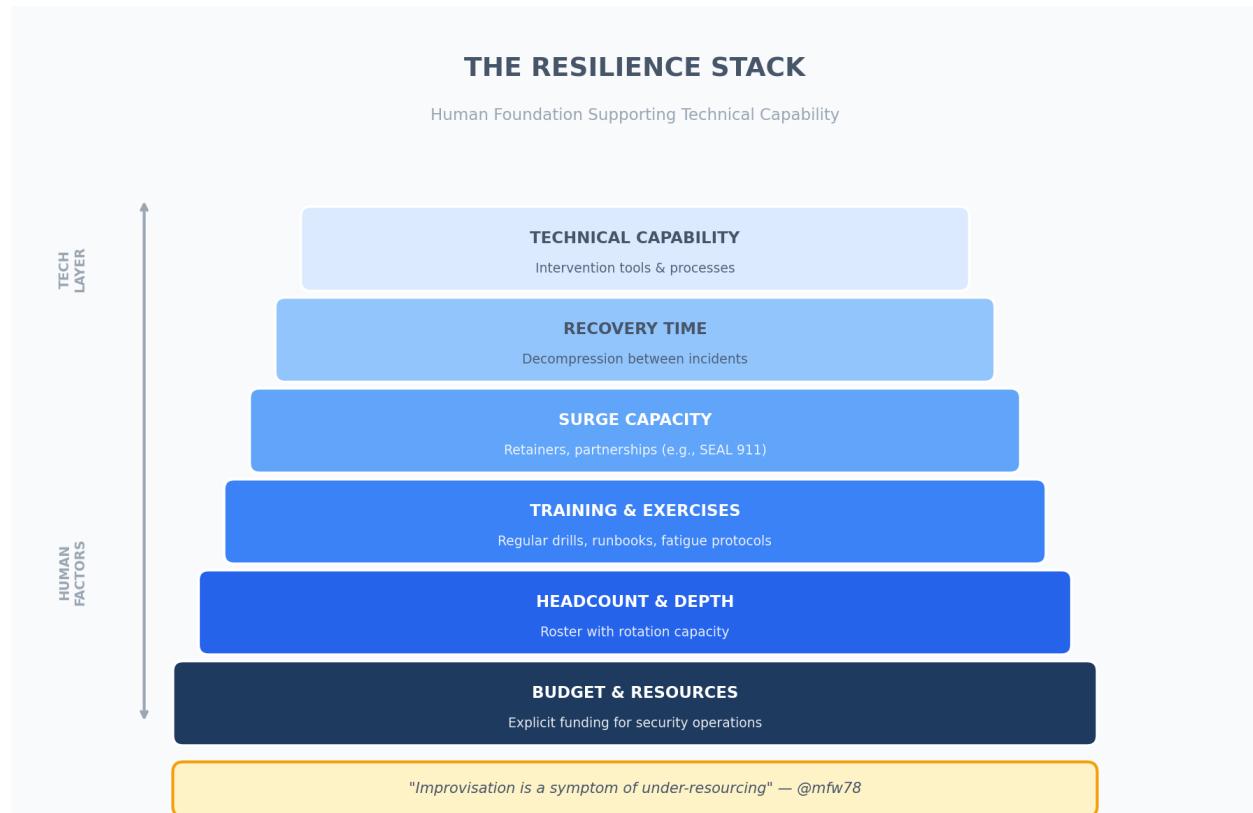


Figure 5: The Resilience Stack shows human factors as the foundation supporting technical capability.

0.7.2 High-Reliability Industry Principles

Factor	Requirement
Headcount	Sufficient depth to allow rotation during extended operations
Surge Capacity	Retainers, partnerships, or standby arrangements
Training	Regular exercises, documented runbooks, fatigue recognition
Recovery Time	Teams need decompression between incidents
Just Culture	Distinguish systemic failures from individual culpability

0.7.3 The Accountant's Question

If GnosisDAO wishes to maintain intervention capability, the cost should be budgeted explicitly:

- Internal headcount with appropriate depth
- External retainers or partnerships with security firms
- Comprehensive bug bounty programmes
- Training and exercises budget
- Recovery capacity: Time and resources for post-incident decompression

Treating emergency response as an externality creates fragility. Explicit budgeting enables informed resourcing decisions.

0.7.4 Integration with the Framework

The proposed Emergency Council should: 1. Have defined roster depth (not just 3 people) 2. Include retainer agreements with external security firms (e.g., SEAL 911) 3. Conduct quarterly exercises (tabletop simulations) 4. Maintain explicit “fatigue protocols” for extended incidents 5. Conduct post-incident “just culture” reviews

0.8 Part 6: Alignment with Gnosis Values

Gnosis Value	How This Framework Supports It
Credible Neutrality	Rules are defined <i>before</i> incidents, not during them
Transparency	All criteria, decisions, and post-mortems are public
Decentralisation	Emergency Council proposes; DAO decides; nodes execute
User Protection	Rapid response capability reduces loss exposure
Sustainability	Explicit resourcing prevents burnout and improvisation

0.9 Proposed Next Steps

0.9.1 For GnosisDAO

1. Review this framework and provide feedback (this forum thread)
2. Integrate all or part of this proposed framework into the working framework by the current working group to refine criteria and procedures
3. Draft a GIP formalising the Intervention Framework
4. Establish Emergency Council composition and authority

0.10 Data Sources & References

Source	Description	Link
Charoenwong & Bernardi (2022)	\$88B in losses (2011–2021)	SSRN 3944435
Charoenwong et al. (2025)	A Framework for Evaluating Regulatory Approaches	SSRN 5368708
Bybit Security Lab (Nov 2025)	166 chains analysed	Report
Anthropic Red Team (Dec 2025)	AI exploit capabilities	red.anthropic.com
LIF Repository	Research repository	GitHub
LIF Technical Extension	Emergency Council power scope (Gnosis-specific)	Technical Extension

0.11 Appendix: Glossary

Term	Definition
Ex Ante	Prevention before execution (vs. Ex Post = after)
Optimistic Freeze	Immediate pause, reverted if not ratified by DAO
Hierarchy of Precision	Graduated intervention levels (Network ⊑ Account)
Emergency Council	Designated multisig with authority to trigger pause
Credible Neutrality	Rules that do not favour any party, defined in advance

Contact: @e3o8o (Forum) | [X/Twitter](#) | [GitHub](#)